

# Kapish Girish Kela

+91 8308199141 | kapishh17@gmail.com | linkedin.com/in/kapish-kela | github.com/Kapish17

## PROFESSIONAL SUMMARY

Computer Science undergraduate (CGPA 8.39/10) with hands-on experience in Machine Learning, Generative AI, and full-stack development across 2 internships and 3 end-to-end projects. Built full-stack ML automation and RAG platforms using FastAPI, LangChain, Google Gemini, FAISS, and BM25. Strong fundamentals in DSA, DBMS, OS, and Computer Networks; proficient in Python, SQL, and Prompt Engineering.

## EDUCATION

**VIT Bhopal University — B.Tech in Computer Science and Engineering (AIML)**

2023 – Present

CGPA: 8.39/10

## EXPERIENCE

**Software Development Intern — HEDGE5, Bangalore**

May 2025 – Aug 2025

- Built a Biogas Calculator tool for sustainability initiatives, automating manual estimation steps and cutting calculation time by roughly 35%.
- Designed and deployed the company website across 5+ pages using HTML, CSS, and JavaScript, with the site indexed on Google within 2 weeks of launch; contributed to agile sprints and 2 deployment cycles over the internship.

**Data Analyst Intern — Bluestock Fintech**

Jun 2026 – Aug 2026

- Cleaned and validated mutual fund datasets from 10+ sources using Python, Pandas, and NumPy, improving downstream data quality and consistency.
- Designed a SQLite database with 6 fact/dimension tables (managed via SQLAlchemy) and wrote 10+ SQL queries to analyze NAV trends, AUM distribution, SIP inflows, and transaction patterns.
- Built a 4-stage analytics workflow (cleaning, validation, DB design, reporting) that reduced manual reporting effort by roughly 30%.

## PROJECTS

**ML-Copilot | Full-Stack ML Automation & Explainability Platform** | *FastAPI, Next.js/React/TypeScript, scikit-learn, Docker, SHAP, Google Gemini* GitHub | Demo

- Built ML-Copilot, a full-stack ML automation platform (FastAPI + Next.js/React/TypeScript) that ingests, profiles, trains and cross-validates multiple candidate models (scikit-learn), and evaluates the winner on a held-out test set; containerized with Docker, deployed as separate services on Render.
- Implemented model explainability with SHAP (permutation-importance fallback), generating global and per-prediction local explanations surfaced through the API and agent tooling; added a RAG layer using a local index and hashing embeddings (no external vector DB) to ground answers in docs/experiment history with citation checking.
- Built an LLM-agnostic AI agent (planner + tool-calling over 4 tools: profiling, experiment execution, explanation, knowledge search) on a provider-abstracted LLM layer (Gemini), with bounded iterations, retry/rate-limit handling, and concurrency throttling for reliable free-tier API usage.

**Multiple Disease Prediction & Healthcare Recommendation System** | *Python, Scikit-learn, Streamlit*

GitHub | Demo

- Trained and unified 3 ML models for Diabetes (~80% accuracy), Heart Disease (~83% accuracy), and Parkinson's (~87% accuracy) prediction on a single platform.
- Combined ML predictions with symptom-based logic using `predict_proba()` to generate probability-based risk assessments across all 3 disease modules.
- Added supporting features including nearby-hospital lookup, personalized lifestyle guidance, and risk-aware recommendations based on predicted probability.

**LocalDocs AI Assistant | RAG-Based Document Q&A System** | *React + Vite, FastAPI, FAISS, BM25, EasyOCR, PyMuPDF, Docker* GitHub | Demo

- Built LocalDocs AI Assistant, a full-stack RAG application (React + Vite, FastAPI) enabling grounded, cited Q&A over user-uploaded documents (PDF/DOCX/PPTX/CSV) via hybrid retrieval blending FAISS vector similarity with BM25 keyword search.
- Engineered per-session document isolation using a custom X-Session-Id header, giving each browser session its own FAISS/BM25 index server-side, with per-page OCR fallback (EasyOCR + PyMuPDF) for scanned PDFs without a text layer.
- Diagnosed and fixed a production memory bottleneck on Render's free tier by migrating embedding generation from a locally-loaded Hugging Face model (~850MB RAM) to a remote inference API call, cutting backend memory usage by over 75% while containerizing the full stack via Docker.

## TECHNICAL SKILLS

**Core CS Subjects:** Data Structures & Algorithms, DBMS, Operating Systems, Computer Networks, Object-Oriented Programming

**Generative AI & LLM:** LangChain, Google Gemini, RAG, Prompt Engineering, FAISS, BM25, SHAP, AI Agents

**Languages:** C++, Python, TypeScript, SQL **Web:** React, Next.js, FastAPI, Streamlit

**ML & Data:** Scikit-learn, Pandas, NumPy, Feature Engineering **Databases & Tools:** MongoDB, SQLite, SQLAlchemy, Docker, Git, GitHub

## CERTIFICATIONS

**Applied Machine Learning** — Coursera | **Cloud Computing** — NPTEL, IIT Kharagpur

**Oracle Cloud Infrastructure 2025 Certified AI Foundations Associate**

## LEADERSHIP & ACTIVITIES

**Cricket Team Captain** — Led team strategy and coordination for inter-college tournaments.

**TCS Recruitment Volunteer** — Assisted in campus recruitment logistics and candidate coordination for 500+ applicants.